Context-Sensitive Synchronization-Sensitive Analysis Is Undecidable

G. RAMALINGAM

IBM T.J. Watson Research Center

Static program analysis is concerned with the computation of approximations of the runtime behavior of programs. Precise information about a program's runtime behavior is, in general, uncomputable for various different reasons, and each reason may necessitate making certain approximations in the information computed. This article illustrates one source of difficulty in static analysis of concurrent programs. Specifically, the article shows that an analysis that is simultaneously both context-sensitive and synchronization-sensitive (that is, a context-sensitive analysis that precisely takes into account the constraints on execution order imposed by the synchronization statements in the program) is impossible even for the simplest of analysis problems.

1. INTRODUCTION

This article illustrates one source of difficulty in statically analyzing concurrent programs. In the realm of static interprocedural analysis, one talks of context-sensitive analyses [Sharir and Pneuli 1981]: these are analyses that are precise with respect to calling-context, in a formally defined way. In a concurrent program that contains synchronization constructs, one can similarly define the notion of synchronization-sensitive analyses [Taylor 1983]: these are analyses that are precise with respect to the synchronization structure of the program. Efficient context-sensitive analysis algorithms have been developed for a number of analysis problems (e.g., see Reps et al. [1995] and Sagiv et al. [1996]). Taylor [1983] shows that intraprocedural synchronization-sensitive analysis of concurrent programs is NP-hard. Consequently, intraprocedural synchronization-sensitive analyses can be expensive, but it is straightforward to show that they are possible for a number of analysis problems. In this article we address the problem of interprocedural analysis of concurrent programs. Specifically, we show that an analysis that is simultaneously both

context-sensitive and synchronization-sensitive is impossible even for the simplest of analysis problems.

2.   THE PROBLEM

In this section we present a simple description of the problem we are addressing. A more formal definition of the problem appears in the Appendix.

We are interested in the static analysis of concurrent programs. We use the programming model described by Taylor [1983], and used by Naumovich and Avrunin [1998]. The model is a simple one, where a concurrent program consists of a set of tasks, which is statically specified. (Thus, new tasks cannot be dynamically created during program execution.) A task is essentially like a sequential (nonconcurrent) program, consisting of a collection of (possibly recursive) procedures, with a distinguished *main* procedure. (Taylor [1983] focuses on the special case where each task consists of a single nonrecursive procedure.) The only concurrency primitive we consider is that of (rendezvous style) synchronization: a task T may *invoke* an *entry* $S.E$ belonging to a task S, and a task T may *accept* an entry belonging to T. Note that the entry invoked or accepted is statically specified. (Thus, entries are *not* first-class values.) This model corresponds to a simple subset of Ada. (Most concurrent progamming languages provide for several more powerful concurrency primitives. The undecidability result established in this article applies to any language that includes at least the above features, or other equivalent features.)

The semantics of the synchronization primitives are as follows. A task which invokes an entry $S.E$ is suspended until the task S *accepts* the entry E, after which both tasks proceed with their execution. If the task S performs an *accept* of entry E when no other task has invoked the entry $S.E$, it is suspended until some other task invokes the entry E, after which both tasks are free to proceed with their execution.

An execution of a concurrent program consists of the execution of each of its tasks (i.e., the *main* procedure of each of its tasks). The execution of each task proceeds sequentially. Except for the constraints imposed by the semantics of the synchronization primitives, statements belonging to different tasks may be executed in any order.

The goal of various static analyses is to compute information about the possible behavior of the program at runtime. The nondeterminism inherent in concurrent programs can complicate static analysis. If the order in which statements belonging to different tasks are executed can affect the program behavior and the values computed, then static analysis has to conservatively account for all possible orders in which statements of different tasks might execute. Since the synchronization primitives impose some constraints on the order in which different statements may execute, static analyses that take these constraints into account can compute more precise results than those do not.

The example in Figure 1 illustrates this. In an execution of Figure 1, statements [2] and [5] will synchronize with each other. Similarly, the execution of statements [3] and [7] will be synchronized. Consequently, the execution of statement [1] is constrained to occur before the execution of statement [6], which is constrained to occur before the execution of statement [4]. Hence, the value of variable V in statement [4] will be the value 1 assigned in statement [6]. Static analysis that

```
task body X is begin              task body Y is begin
   [1] V := 0;                       [5] accept E;
   [2] invoke Y.E;                   [6] V := 1;
   [3] invoke Y.E;                   [7] accept E;
   [4] Z := V;                    end Y;
end X;
```

Fig. 1.   Example of rendezvous style synchronization (with Ada-like syntax).

takes the synchronization constraints into account can infer this. An analysis of this program that ignores the synchronization constraints, however, will have to assume that the statements [1] and [6] may execute in any order, and, hence, yield the less precise information that the value of variable V in statement [4] might be either 0 or 1.

We say that a static analysis is synchronization-sensitive if it precisely takes into account the constraints imposed by the synchronization statements in the program. A more formal definition of this concept appears in the Appendix. The (partial) execution of a concurrent program can be represented by a collection of (partial) execution paths, one for each task in the program. We refer to such a collection of execution paths as a *potential execution sequence* for the program. An execution path in a task contains a sequence of *synchronization events*, where a synchronization event is either the invocation or acceptance of an entry. Not every potential execution sequence can occur in practice. In particular, the semantics of synchronization rules out some potential execution sequences as being infeasible. We say that a given potential execution sequence is a legal execution sequence if the synchronization events in the sequence can be paired up in a consistent manner. (We refer the reader to the Appendix for a more detailed and formal definition of these concepts.) A static analysis is said to be synchronization-sensitive if it treats potential execution sequences that are not legal as unexecutable and computes information describing only execution along legal execution sequences (similar to the meet-over-all-paths analyses [Kildall 1973] and meet-over-all-valid-paths analyses [Sharir and Pneuli 1981]).

Synchronization-sensitive analysis is not easy. The synchronization constraints can impose not only an order on the execution of statements belonging to different tasks, it can also rule out certain paths within a single task as being unexecutable. This is illustrated in Figure 1.

In Figure 1, the invocation of entry E in statement [1] is accepted by statement [9]. If task X takes the true branch of the conditional in line [2], then the execution of [3] will suspend task X forever, since the entry invocation will not be accepted by task Y. Hence, if execution reaches statement [8] then the false branch *must* have been taken in line [2]. Hence, a *synchronization-sensitive* analysis of the above program can establish that the value of W in line [8] must be 1. The above example illustrates the "power" of synchronization-sensitive analysis, and it should come as no surprise that such analyses can be expensive. Taylor [1983] shows that various such analyses are NP-complete in the intraprocedural case, where each task consists of a single nonrecursive procedure.

Though potentially expensive, such analyses are possible in the intraprocedural

```
task body X is begin
   [1]invoke Y.E;
   [2]if  (···) then
   [3]   invoke Y.E;
   [4]   W := 0;
   [5]else
   [6]   W := 1;
   [7]end if;
   [8]Z := W;
end X;
```

```
task body Y is begin
   [9] accept E;
end Y;
```

Fig. 2. Example showing that synchronization constraints can rule out certain paths within a single task as being unexecutable.

case. Assume that each task is represented by a control-flow graph. Note that for any such multitasking program we can create a "product graph" of the control-flow graphs of the tasks in the program. A vertex in the product graph corresponds to a tuple consisting of a vertex in every task's graph. We have some choices in terms of how the set of edges of the product graph is defined. Consider the product graph that contains an edge from a tuple $(u_1, \cdots, u_k)$ to a tuple $(v_1, \cdots, v_k)$ if and only if every v_i is either equal to u_i or is a successor of u_i and at least one v_i is a successor of u_i. Paths in this product graph represent essentially the set of all potential execution sequences through the given program. On the other hand, it is straightforward to define the set of edges in the product graph such that paths in the product graph represent essentially the set of all legal or synchronization-sensitive execution sequences through the given program. We only need to eliminate those edges that cause a synchronizing vertex u_i to transition to its successor vertex v_i without a matching synchronizing vertex making its transition. We can use this product graph to directly adapt intraprocedural analyses of sequential programs to analyze concurrent programs in a synchronization-sensitive way.

The focus of this article is on the interprocedural analysis of concurrent programs, where each task consists of a set of (possibly recursive) procedures

How do *procedures* complicate such synchronization-sensitive analysis? In particular, assume that each task consists of a set of (possibly recursive) procedures, with a distinguished *main* procedure. Is it feasible to perform *interprocedural* static analyses of concurrent programs in a *context-sensitive and synchronization-sensitive* way? (An interprocedural analysis is said to be context-sensitive if it considers only paths with no mismatching call and return edges. A more complete definition appears in the Appendix. We say that a potential execution sequence is context-sensitive if every path in the sequence is context-sensitive.)

One of the simplest static analysis problem is to determine if program execution might ever reach a particular vertex. We define this *reachability* problem as follows.

Definition 1 (*The Context-Sensitive and Synchronization-Sensitive Reachability Problem*). Given a program P consisting of a set T of tasks, a task τ in T, and a program point u in τ, determine if there exists a context-sensitive legal execution sequence $(\pi_{\tau'} : \tau' \in T)$ for P, where every path $\pi_{\tau'}$ starts at the start vertex of the main procedure of τ' and the path π_τ ends at program point u.

Taylor defines several other analysis problems for concurrent programs. The goal of the "Possible Rendezvous" (or PR) problem is to determine if two given statements in a concurrent program may possibly rendezvous (i.e., synchronize with each other). Another problem closely related to the PR problem is the "May Happen in Parallel" (or MHP) problem, which is to determine if two given statements in a concurrent program may ever execute in parallel. We can, in an analogous fashion, also define generalizations of various analysis problems that have been studied in the context of sequential programs, such as the "reaching-definitions" problem.

The context-sensitive and synchronization-sensitive reachability problem defined above is simpler than the context-sensitive and synchronization-sensitive versions of most analysis problems, including the PR problem, the MHP problem, and the reaching-definitions problem. In other words, the reachability problem can be trivially reduced to these other problems. Consequently, the undecidability of the reachability problem (established in the next section) implies the undecidability of these other analyses.

3. AN UNDECIDABILITY RESULT

This section shows that identifying the existence of context-sensitive synchronization-sensitive paths is impossible, implying the impossibility of performing static analysis of concurrent programs in a context-sensitive and synchronization-sensitive way.

3.1 A Reduction From Post's Correspondence Problem

Definition 2. The Post's Correspondence Problem or PCP is the following: given arbitrary lists A and B of r strings each in $\{0,1\}^+$, say

$$A = w_1, w_2, \cdots, w_r$$
$$B = z_1, z_2, \cdots, z_r$$

does there exist an integer k greater than zero and a sequence of integers i_1, i_2, $\cdots$, i_k such that $w_{i_1} w_{i_2} \cdots w_{i_k} = z_{i_1} z_{i_2} \cdots z_{i_k}$?

THEOREM 1. *The Post's Correspondence Problem is undecidable [Hopcroft and Ullman 1979].*

Note that each w_i and z_j in the above definition is a string of bits, not a single bit. Further, note that k does not denote a fixed value in the definition of the problem. The problem also does allow repetitions in the sequence of integers i_1, i_2, $\cdots$, i_k. These are the factors that make the problem undecidable.

THEOREM 2. *The context-sensitive synchronization-sensitive reachability problem is undecidable.*

PROOF. We will show a reduction from the Post's Correspondence Problem to the context-sensitive synchronization-sensitive reachability problem. In particular, given an instance

$$A = w_1, w_2, \cdots, w_r$$
$$B = z_1, z_2, \cdots, z_r$$

```
task D is
  entry allDone_A, allDone_B;
end D;

task A is
  entry W_1, W_2, ···, W_r;
end A;

task B is
  entry Z_1, Z_2, ···, Z_r;
  entry bit_0, bit_1;
end B;
```

```
task body A is
  procedure F is
    if  (···) then
      return ;
    elsif  (···) then
      accept W_1;
      F();
      // Encode string w_1 in reverse.
      // Note: w_1[|w_1|] denotes the last bit of w_1.
      invoke B.bit_{w_1[|w_1|]}; ···; invoke B.bit_{w_1[1]};
    elsif  (···) then
    ···
    else
      accept W_r;
      F();
      invoke B.bit_{w_r[|w_r|]}; ···; invoke B.bit_{w_r[1]};
    end if;
  end F;
begin
  F();
  invoke D.allDone_A;
end A;
```

```
task body D is begin
  repeat
    if  (···) then
      invoke A.W_1; invoke B.Z_1;
    elsif  (···) then
      invoke A.W_2; invoke B.Z_2;
    ···
    else
      invoke A.W_r; invoke B.Z_r;
    end if;
  until  (···);
  accept allDone_A;
  accept allDone_B;
  // The following statement is
  // reachable iff the given PCP
  // instance has a solution
  u:
end D;
```

```
task body B is
  procedure H is
    if  (···) then
      return ;
    elsif  (···) then
      accept Z_1;
      H();
      // Encode string z_1 in reverse.
      // Note: z_1[|z_1|] denotes the last bit of z_1.
      accept bit_{z_1[|z_1|]}; ···; accept bit_{z_1[1]};
    elsif  (···) then
    ···
    else
      accept Z_r;
      H();
      accept bit_{z_r[|z_r|]}; ···; accept bit_{z_r[1]};
    end if;
  end F;
begin
  H();
  invoke D.allDone_B;
end B;
```

Fig. 3.   An instance of the reachability problem generated from an instance of Post's Correspondence Problem.   The exit vertex of task D is reachable along a context-sensitive and synchronization-sensitive execution sequence if and only if the given PCP instance has a solution.

of PCP, consider the multitasking program (written using syntax similar to that of Ada) shown in Figure 3.

This program consists of three tasks D, A, and B. We can establish a bijective correspondence between complete execution paths in task D and nonempty sequences of integers i_1, i_2, $\cdots$, i_k, where every i_j lies between 1 and r (inclusive). In particular, any execution path that iterates through the repeat-until loop k times corresponds to a sequence i_1, i_2, $\cdots$, i_k where i_j indicates which of the r branches inside the loop was taken during the jth iteration of the loop. One can think of task D as "generating" all possible integer sequences (of the above form) and "communicating" the generated sequence to tasks A and B.

Task A consists of a recursive procedure F. Procedure F "receives" the sequence generated by task D, and utilizes the recursive call-chain to "remember" this sequence. Once a sequence i_1, i_2, $\cdots$, i_k is received, procedure F "generates" the string $w_{i_1} w_{i_2} \cdots w_{i_k}$ in reverse and "communicates" this generated string to task B.

Task B also consists of a recursive procedure H. Procedure H is very similar to procedure F. It "receives" the sequence generated by task D, and utilizes the recursive call-chain to "remember" this sequence. It then determines the string $z_{i_1} z_{i_2} \cdots z_{i_k}$ in reverse and is ready to "accept" this string from task A.

If $w_{i_1} w_{i_2} \cdots w_{i_k}$ is a suffix of $z_{i_1} z_{i_2} \cdots z_{i_k}$, then procedure F will terminate, and task A will signal task D that it is all done by invoking entry $D.allDone_A$. Similarly, B will invoke entry $D.allDone_B$ if $z_{i_1} z_{i_2} \cdots z_{i_k}$ is a suffix of $w_{i_1} w_{i_2} \cdots w_{i_k}$. In particular, A will invoke $D.allDone_A$ and B will invoke $D.allDone_B$ if and only if $w_{i_1} w_{i_2} \cdots w_{i_k}$ is equal to $z_{i_1} z_{i_2} \cdots z_{i_k}$.

Thus, if the given PCP instance has a solution, then the exit vertex (the statement labeled u) of task D is reachable along a context-sensitive synchronization-sensitive execution sequence. Conversely, if the exit vertex of task D is reachable along a context-sensitive synchronization-sensitive execution sequence, then the given PCP instance has a solution.

It follows that the context-sensitive synchronization-sensitive reachability problem is undecidable.   □

3.2   An Alternative Proof

We briefly mention an alternative reduction that can be used to establish the above undecidability result. Not all undecidability proofs are equivalent, as each proof may offer different insights into the kind of approximations and simplifications that can be used to tackle an undecidable problem. It is known that it is undecidable if the languages defined by two arbitrary context-free grammars have a nonempty intersection [Hopcroft and Ullman 1979, Theorem 8.10]. Given a context-free grammar G_1, the grammar can be encoded as a process as follows: represent each terminal of the grammar by an entry; represent each nonterminal of the grammar by a procedure consisting of a multiway conditional statement consisting of a branch for every production of that nonterminal. The branch encodes the right-hand side of the corresponding production, encoding the occurrence of a terminal by an invocation of the corresponding entry, and encoding the occurrence of a nonterminal by a call to the corresponding procedure. Execution of any (context-sensitive) path from the start vertex to the exit vertex of a procedure results in a sequence of entry invocations that "belongs" to the language defined by the nonterminal cor-

responding to the procedure. Conversely, for any string in the language defined by the nonterminal, there exists a (context-sensitive) path from the start vertex to the exit vertex of the procedure whose execution results in the corresponding sequence of entry invocations. Thus, the "entry invocations" performed by the procedures of the above process correspond to the languages defined by the nonterminals of G_1. Dually, one can encode a context-free grammar G_2 by a system of procedures that perform a sequence of "entry acceptances" corresponding to the languages defined by the nonterminals of G_2. Thus, it is straightforward to define a program consisting of two processes that will terminate successfully if and only if the languages defined by the two given context-free grammars have a nonempty intersection.

4. DISCUSSION

It is well known that "precise static program analysis" is impossible (e.g., through connections to the halting problem). The result established in this article, however, says more than that, and we briefly elaborate on this issue here.

All static program analysis algorithms combat undecidability by working with a *nonstandard* semantics of programs, which serves as an abstraction or approximation of the standard semantics of programs [Kildall 1973; Cousot and Cousot 1977]. Given a nonstandard semantics, one can talk of the most precise analysis with respect to that semantics, referred to as the collecting interpretation with respect to the nonstandard semantics, or the meet-over-all-paths solution [Kildall 1973]. (We are being somewhat loose with our terminology here for the sake of brevity.) One question that arises in designing static analysis algorithms is what kind of approximations one must make in the nonstandard semantics to guarantee that the most precise analysis with respect to that semantics is computable.

Consider procedureless sequential programs. If the set of program states in the nonstandard semantics is finite (or, in Kildall's formalism, if the "lattice" is finite), then the most precise analysis with respect to the semantics is computable (using iterative techniques, for example).

In the presence of (potentially recursive) procedures, the program's state, in the standard semantics, consists of two components, one of which is the "call stack." We will refer to the second component of the program state as the "global state," for the lack of a better term. (Note that the "call stack" is only a sequence of procedures; if the procedures have local variables, we consider them part of the "global state.") One can construct nonstandard semantics that approximate the global state without approximating the call stack. The set of program states in such a nonstandard semantics is infinite, since the call stack can grow unbounded. Interestingly, the most precise analysis with respect to such a nonstandard semantics is still computable as long as the set of global states in the nonstandard semantics is finite (as shown by Sharir and Pneuli [1981]). Such analyses are precise with respect to the calling-context and are referred to as context-sensitive analyses.

Now consider concurrent programs without procedures. One can construct nonstandard semantics for such programs by approximating the state of each task, but preserving the synchronization semantics. The most precise analysis with respect to such a nonstandard semantics is computable if the set of abstract states for each task is finite. This leads to the notion of synchronization-sensitive analysis.

Unfortunately, these positive results fail to hold if we consider static analysis of concurrent programs with recursive procedures. The result in this article shows that if a nonstandard semantics approximates neither the synchronization semantics of the program nor the call stack, then the most precise analysis with respect to that semantics is not computable.

In terms of related work, Reif [1979] considers the problem of intraprocedural analysis of a system of communicating processes. In Reif's model, processes transmit and receive messages through asynchronous communication channels. Reif shows that reachability analysis is undecidable in the case of *dynamic* communication systems, where channels are first-class values, and that reachability analysis is EXP-SPACE hard in the case of *static* communication, where the channels specified in "receive" and "transmit" statements are constants. Our result shows that the problem is undecidable even in the case of static communication in the interprocedural case. The model we use is essentially an interprocedural extension of the model Taylor [1983] uses. Taylor shows that several (intraprocedural) analysis problems are NP-complete in his model. On a more positive note, Esparza and Podelski [2000] present efficient algorithms for "context-sensitive" interprocedural analysis of parallel programs built using parbegin-parend constructs (which may be thought of as "structured" synchronization primitives). Even though these algorithms do not address arbitrary synchronization statements, the framework used by Esparza and Podelski, based on process algebra, itself appears to be a convenient one for formalizing and investigating analysis of parallel programs containing arbitrary synchronization primitives as well.

5. CONCLUSION

The precision of static analysis of concurrent programs, especially programs where different tasks (threads) may access and modify "shared" variables, can obviously be improved by taking into account the sequencing constraints imposed on the statements of the different tasks (threads) by the concurrency primitives. Several people (e.g., see Duesterwald and Soffa [1991], Masticola and Ryder [1993], Grunwald and Srinivasan [1993], Masticola [1993], Naumovich and Avrunin [1998], and Naumovich et al. [1998]) have proposed various conservative algorithms for identifying the sequencing constraints on the statements of a concurrent program. This article shows that identifying such sequencing constraints in a context-sensitive way is impossible. The proof technique used in this article was inspired by Reps' result [Reps 1999] which shows that context-sensitive, structure-transmitted data-dependence analysis is impossible. The actual undecidability result may not be very surprising once one understands the problem, but we believe that it is useful to understand these limitations when designing static analyses. We are not aware of any equivalent result in the literature, despite the importance of this problem and the amount of work that has gone into designing (conservative) static analyses for concurrent programs.

APPENDIX

We present a formal definition of the problem considered in the article below. The definition extends the traditional concepts of meet-over-all-paths analysis and context-sensitive analysis to define the concept of synchronization-sensitive analysis

analogously. The definition is based on the model introduced by Taylor [1983] and used by several others. This formalism is based on a control-flow graph representation of procedures, but can easily be adapted for other representations.

Sequential Single-Procedure Programs

We will first consider sequential (i.e., nonconcurrent) programs. In the absense of procedures and procedure-calls, a program can be represented by a simple control-flow graph. A single-procedure control-flow graph G is a tuple $\langle V, E, s, e \rangle$ consisting of a set of vertices V, a set of edges $E \subseteq V \times V$, a special "start" vertex s, and a special "exit" vertex e. We assume, without loss of generality, that the start vertex s has no predecessors and that the exit vertex e has no successors. We say that (u, v) is an edge in the graph if $(u, v) \in E$. A sequence of vertices $[u_1, u_1 \cdots, u_k]$ is said to be a path in the graph if, for $1 \leq i \leq k$, (u_i, u_{i+1}) is an edge in the graph. We will denote the set of vertices of a single-procedure control-flow graph G by $V(G)$.

Sequential Multiprocedure Programs

Now consider a (sequential) program consisting of several procedures. We can represent such a program by a collection of single-procedure control-flow graphs, one for every procedure in the program, where some vertices are distinguished as being "call sites." A call site vertex denotes a procedure call, and for such vertices we assume that the called procedure is known.

More formally, a multiprocedure control-flow graph MPG is defined to be a tuple $\langle P, \langle G_p : p \in P \rangle, main, callee \rangle$ consisting of a set of procedures P, a distinguished procedure $main$ belonging to P, a single-procedure control-flow graph G_p for every procedure p in P, as well as a partial function $callee$ from the set $\cup_{p \in P} V(G_p)$ to P. We refer to $\cup_{p \in P} V(G_p)$ as the set of vertices of MPG and denote it by $V(MPG)$. A vertex u in MPG is said to be a callsite if $callee$ is defined for that vertex, in which case we refer to $callee(u)$ as the procedure called at call site u. We assume, for the sake of simplicity, that callsites are distinct from the start or exit vertices of any single-procedure control-flow graph. We also assume, for the sake of simplicity, that a callsite has exactly one successor. It is straightforward to ensure that these assumptions are satisfied, by introducing some dummy vertices if necessary. We will refer to the start and exit vertex of the $main$ procedure as the start and exit vertex of the program.

Edges in a multiprocedure graph can be classified into two kinds. Edges whose source is not a callsite are *intraprocedural* edges. Edges whose source is a callsite are *return-fall-through* edges. A frequently used alternative representation of a multiprocedure graph is the "supergraph," which is obtained by taking the union of the graphs G_p and replacing every return-fall-through edge $u \rightarrow v$ by a pair of labeled edges $u \rightarrow s(G_{callee(u)})$ (an edge from the call site to the start vertex of the called procedure, referred to as a "call edge") and $e(G_{callee(u)}) \rightarrow v$ (an edge from the exit vertex of the called procedure to the successor of the call site, referred to as a "return edge"). We attach a label "$(_u$" to the call edge and a label "$)_u$" to the return edge. (Recall that u denotes the callsite.) The labeling on the edges serves to match call edges with the corresponding return edges. A sequence of vertices

in a multiprocedure graph is said to be a *supergraph path* if it forms a path in the corresponding supergraph.

Intraprocedural analysis algorithms (i.e., algorithms designed for procedureless programs) can be adapted to analyze programs with (potentially recursive) procedures by treating a supergraph as a single-procedure control-flow graph. Such algorithms are *context-insensitive* interprocedural analysis algorithms. These algorithms are conservatively safe, since every execution of a multiprocedure program corresponds to a path in the supergraph.

However, not every path in the supergraph corresponds to an execution of the multiprocedure program. One property possessed by a path corresponding to an execution of the program is that it does not contain any mismatched pair of call and return edge. More formally, a supergraph path is said to be a *same-level valid path* if the sequence of labels on the call and return edges of the path belongs to the language of balanced parentheses generated from the nonterminal *matched* by the following context-free grammar:

$$matched \rightarrow (_u \; matched \;)_u \; matched \quad \text{for every call site } u$$
$$| \; \epsilon$$

A supergraph path is said to be an *(interprocedurally) valid path* if the sequence of labels on the call and return edges of the path belongs to the language generated from the nonterminal *valid* by the following context-free grammar:

$$valid \rightarrow valid \, (_u \; matched \quad \text{for every call site } u$$
$$| \; matched$$

The precision of interprocedural analysis can be improved by exploiting the fact that only valid paths that are executable. Algorithms that treat the set of interprocedurally valid paths as an approximation to the set of all executable paths and compute the "meet-over-all-valid-paths" solution [Sharir and Pneuli 1981] are referred as *context-sensitive* algorithms. An interprocedurally valid path is also referred to as a *context-sensitive* path sometimes.

Concurrent Programs

We now turn our attention to multitasking concurrent programs. We use the model described by Taylor [1983], and used by Naumovich and Avrunin [1998]. The programming model is a simple one, where the set of tasks is statically specified. A multitasking program consists of a collection of tasks, each task being represented by a multiprocedure graph. The only concurrency primitive we consider is that of synchronization: a task T may invoke an *entry* $S.E$ belonging to a specific task S, in which case the task T waits until the task S *accepts* the entry E. If the task S performs an *accept* of entry E when no other task has invoked the entry $S.E$, it waits until some task invokes the entry. Note that the entry invoked or accepted is statically specified. (Thus, entries are *not* first-class values.) This model corresponds to a simple subset of Ada. (Most concurrent progamming languages provide for several more powerful concurrency primitives. Our undecidability result applies to any language that contains the above features or primitives that can be used to simulate the above features directly.)

We can formalize these ideas as follows. A multitasking program is a tuple $\langle T, E, G, accepts, invokes \rangle$, consisting of a set of tasks T, a multiprocedure graph G_τ for every task τ in T, a set of entry names E, a partial function *accepts* from the set of vertices $\cup_{\tau \in T} V(G_\tau)$ to E, and a partial function *invokes* from the set of vertices $\cup_{\tau \in T} V(G_\tau)$ to $T \times E$. Further, for any vertex at most one of *accepts* or *invokes* is defined. A vertex for which *invokes* is defined is said to be an entry invocation vertex, while a vertex for which *accepts* is defined is said to be an entry acceptance vertex. (Note that every entry *belongs* to a particular task. We identify an entry by a pair consisting of the task that owns the entry and an entry name. Thus, different entries belonging to different tasks may have the same name. A task may *invoke* an entry belonging to any task, while a task can *accept* only an entry belonging to itself. This is why we choose to model the function *invokes* as one that returns a pair belonging to $T \times E$, while we model the function *accepts* as one returning only an element of E. An *accepted* entry is implicitly assumed to belong to the *accepting* task.)

We now formalize some concepts related to the execution of a multitasking program. An execution of a multitasking program corresponds to the execution of all the tasks in the program. Hence, we may represent an execution of the program by a collection $\Pi = (\pi_\tau : \tau \in T)$, where each π_τ is a supergraph path in G_τ. We refer to such a collection of paths as a *potential execution sequence* through the program. We will refer to the potential execution sequence as being context-sensitive if every supergraph path is an interprocedurally valid path. Further, we will say that the sequence begins at the program's start vertex set if the first vertex of every supergraph path is the corresponding task's start vertex.

A potential execution sequence specifies the order in which vertices in a given task are executed, but specifies nothing about the order in which vertices belonging to different tasks are executed. Not every potential execution sequence is a legal execution sequence. In particular, a legal execution sequence has to respect the rendezvous constraints in the program. We formalize these concepts below.

Note that a vertex in a task may execute multiple times. In other words, a path π_τ may consist of multiple occurrences of the same vertex. We will refer to an *occurrence* of a vertex in a path as an *event*. We will identify an event corresponding to the ith element of a path π_τ by the pair (π_τ, i). (Recall that a path is just a sequence of vertices.) Formally, let π_τ denote some path in the supergraph of a task τ. We define the set of events in π_τ, $Events(\pi_\tau)$, to be the set $\{(\pi_\tau, i) \mid 1 \le i \le |\pi_\tau|\}$, where $|\pi_\tau|$ denotes the length of the sequence π_τ. The set of events in a potential execution sequence $\Pi = (\pi_\tau : \tau \in T)$ is defined to be $\cup_{\tau \in T} Events(\pi_\tau)$ and will be denoted by $Events(\Pi)$.

The sequencing constraints in a multitasking program, induced by the rendezvous constructs, constrain the order in which the different events may occur. An event (π, i) is said to be a *synchronization event* if the vertex $\pi(i)$ is either an acceptance vertex or an invocation vertex. (For any sequence π, we will denote the ith element of the sequence by $\pi(i)$.) We denote the set of all syncronization events in a potential execution sequence Π by $SyncEvents(\Pi)$. A synchronization event (π, i) is said to be a *potential matching event* of another synchronization event (π', i') if the two vertices $\pi(i)$ and $\pi'(i')$ denote an invocation vertex of an entry (τ, e) and an acceptance vertex of the entry e in task τ.

A *potential rendezvous mapping* ξ is a symmetric bijection on the set of all synchronization events of a potential execution sequence that maps every synchronization event to another potential matching event. In other words, ξ is a function from $SyncEvents(\Pi)$ to $SyncEvents(\Pi)$ such that (i) $\xi(e) = e'$ iff $\xi(e') = e$ and (ii) $\xi(e)$ is a potential matching event for e.

For any path π_τ, its *event-ordering graph* $D(\pi_\tau)$ is the graph whose set of vertices is $Events(\pi_\tau)$ and whose set of edges is $\{(\pi_\tau, i) \rightarrow (\pi_\tau, i+1) \mid 1 \leq i < |\pi_\tau|\}$.

Given a potential rendezvous mapping ξ for a potential execution sequence $\Pi = (\pi_\tau : \tau \in T)$, we define the event-ordering graph $D(\Pi, \xi)$ to consist of the union of the $D(\pi_\tau)$ for every τ in T combined with the set of "synchronization" edges ξ. More precisely, $D(\Pi, \xi)$ consists of the set of vertices $Events(\Pi)$ and the set of edges $\{(\pi_\tau, i) \rightarrow (\pi_\tau, i+1) \mid \pi_\tau \in \Pi, 1 \leq i < |\pi_\tau|\} \cup \{ e \rightarrow \xi(e) \mid e \in SyncEvents(\Pi) \}$.

Thus, $D(\Pi, \xi)$ represents the constraints imposed by ξ on the order in which events of Π can occur. These constraints include the total ordering on the events within a single task, as well as the constraint that two events e and $\xi(e)$ that synchronize with each other must occur simultaneously. The synchronization constraint is modeled by two edges, one from e to $\xi(e)$ and another from $\xi(e)$ to $\xi(\xi(e))$ $= e$. Thus, for any synchronization event e, e and $\xi(e)$ form a cycle in the graph $D(\Pi, \xi)$. We will refer to such cycles as *matching two-cycles*. We say that two events e and e' *may happen in parallel* in $D(\Pi, \xi)$ if there exists no path in $D(\Pi, \xi)$ either from e to e' or from e' to e.

A potential rendezvous mapping ξ is said to be *consistent* with the potential execution instance Π if it contains no cycles other than the matching two-cycles.

A potential execution sequence Π is said to be a *synchronization-sensitive* execution sequence if there exists a potential rendezvous mapping consistent with the execution sequence. (Taylor [1983] refers to these sequences as *legal* execution sequences.) Recall that a potential execution sequence is said to be context-sensitive if every supergraph path in the sequence is an interprocedurally valid path.

We now present a more precise and formal version of Definition 1.

Definition 1 (*The Context-Sensitive and Synchronization-Sensitive Reachability Problem*). Given a multitasking program P, a task τ in P, and a vertex u in τ's supergraph, determine if there exists a context-sensitive and synchronization-sensitive execution sequence $(\pi_{\tau'} : \tau' \in T)$ for P, starting from P's start vertex set, in which the last vertex of π_τ is u.

Taylor defines several other analysis problems for concurrent programs. We say that a vertex u may *possibly (context-sensitively) rendezvous* with a vertex v if and only if there exists a (context-sensitive) legal execution sequence Π with a consistent rendezvous mapping ξ and $u = \pi(i)$ and $v = \pi'(i')$ and $\xi((\pi, i)) = (\pi', i')$ for some events (π, i) and (π', i'). The goal of the (context-sensitive) "Possible Rendezvous" (or PR) problem is to determine if two given vertices in a concurrent program may possibly (context-sensitively) rendezvous.

Another problem closely related to the PR problem is the "May Happen in Parallel" (or MHP) problem, which is to determine if two given vertices in a concurrent program may ever execute in parallel. We say that a vertex u may *(context-sensitively) happen in parallel* with a vertex v if and only if there exists a (context-sensitive) legal execution sequence Π with a consistent rendezvous mapping ξ and

$u = \pi(i)$ and $v = \pi'(i')$ for some two events (π, i) and (π', i')) that may happen in parallel.

We can, in an analogous fashion, also define generalizations of various analysis problems (e.g., the "reaching definitions" problem) that have been studied in the context of sequential programs.

REFERENCES

COUSOT, P. AND COUSOT, R. 1977. Abstract interpretation: A unified lattice model for static analysis of programs by construction or approximation of fixed points. In *Conference Record of the Fourth ACM Symposium on Principles of Programming Languages*. Los Angeles, California, 238–252.

DUESTERWALD, E. AND SOFFA, M. L. 1991. Concurrency analysis in the presence of procedures using a data flow framework. In *Proceedings of the ACM SIGSOFT Fourth Workshop on Software Testing, Analysis, and Verification*. 36–48.

ESPARZA, J. AND PODELSKI, A. 2000. Efficient algorithms for pre* and post* on interprocedural flow graphs. In *Conference Record of the 27th ACM Symposium on Principles of Programming Languages*. 1–11.

GRUNWALD, D. AND SRINIVASAN, H. 1993. Efficient computation of precedence information in parallel programs. In *Languages and Compilers for Parallel Computing*. Lecture Notes in Computer Science, vol. 768. Springer-Verlag, 602–616.

HOPCROFT, J. E. AND ULLMAN, J. D. 1979. *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley, Reading, MA.

KILDALL, G. 1973. A unified approach to global program optimization. In *Conference Record of the First ACM Symposium on Principles of Programming Languages*. ACM, New York, NY, 194–206.

MASTICOLA, S. AND RYDER, B. 1993. Non-concurrency analysis. In *Proceedings of the Twelvth Symposium on Principles and Practices of Parallel Programming*. 129–138.

MASTICOLA, S. P. 1993. Static detection of deadlocks in polynomial time. Ph.D. thesis, Rutgers University.

NAUMOVICH, G. AND AVRUNIN, G. S. 1998. A conservative data flow algorithm for detecting all pairs of statements that may happen in parallel. In *Proceedings of ACM SIGSOFT Sixth International Symposium on the Foundations of Software Engineering*. 24–34.

NAUMOVICH, G., AVRUNIN, G. S., AND CLARKE, L. A. 1998. An efficient algorithm for computing MHP information for concurrent Java programs. Tech. Rep. Technical Report 98-44, Department of Computer Science, University of Massachusetts, Amherst, MA.

REIF, J. H. 1979. Data flow analysis of communicating processes. In *Conference Record of the Sixth ACM Symposium on Principles of Programming Languages*. San Antonio, TX, 257–268.

REPS, T. 1999. Undecidability of context-sensitive data-dependence analysis. Tech. Rep. TR-1397, Computer Sciences Department, University of Wisconsin, Madison, WI. Mar. To appear in *ACM Transactions on Programming Languages and Systems*.

REPS, T., HORWITZ, S., AND SAGIV, M. 1995. Precise interprocedural dataflow analysis via graph reachability. In *Conference Record of the Twenty-Second ACM Symposium on Principles of Programming Languages*. 49–61.

SAGIV, M., REPS, T., AND HORWITZ, S. 1996. Precise interprocedural dataflow analysis with applications to constant propagation. *Theoretical Computer Science 167*, 131–170.

SHARIR, M. AND PNEULI, A. 1981. Two approaches to interprocedural data flow analysis. In *Program Flow Analysis: Theory and Applications*, S. Muchnick and N. Jones, Eds. Prentice Hall, Englewood Cliffs, NJ, 189–233.

TAYLOR, R. N. 1983. Complexity of analyzing the synchronization structure of concurrent programs. *Acta Informatica 19*, 57–84.